

# A user guide for *panelstat*

Paulo Guimaraes, BPlim

28jun2018

## 1 Introduction

`panelstat` is a Stata user-written command to help understand the characteristics of standard panel data sets. Usage of `panelstat` command is quite simple:

```
panelstat panelvar timevar [if] [in] [ , options]
```

where `panelvar` is a unit identifier for the panel and `timevar` identifies the time variable. `panelstat` has many options which we will discuss using examples. For now, if you want to see the full list of options simply check the help file for `panelstat`.

## 2 Basic Usage

To illustrate the usage of `panelstat` we start by loading Stata's *nlswork* file set which is a sample of the National Longitudinal Survey of young women aged 14 to 26 years in 1968, and which were observed for several years.

```
. sysuse nlswork.dta, clear
(National Longitudinal Survey. Young Women 14-26 years of age in 1968)
```

The unit identifier for this panel is *idcode* – a unique identifier for each participant in the study. The time variable is *year*. To obtain basic descriptive statistics for this panel we can run `panelstat`:

```
. panelstat idcode year

*****
Analyzing ./nlswork.dta
*****

*****
Basic descriptives
*****
There are 28534 time x individuals observations
There are 4711 unique individuals
Time values range from 68 to 88
```

Maximum time range is 21  
The average number of periods per individual is 6.06  
The level of completeness is 28.84% (100% is a fully balanced panel)  
Average number of gaps per individual is 2.75  
Average gap size is 1.84  
Largest gap is 19

\*\*\*\*\*

\*\*\*\*\*

Distribution of number of observations per individual

\*\*\*\*\*

Observ per individual	Freq.	Percent	Cum.
1	547	11.61	11.61
2	498	10.57	22.18
3	484	10.27	32.46
4	411	8.72	41.18
5	421	8.94	50.12
6	398	8.45	58.57
7	345	7.32	65.89
8	323	6.86	72.74
9	302	6.41	79.16
10	270	5.73	84.89
11	202	4.29	89.17
12	158	3.35	92.53
13	147	3.12	95.65
14	119	2.53	98.17
15	86	1.83	100.00
Total	4,711	100.00	

\*\*\*\*\*

Number of individuals per time unit

\*\*\*\*\*

Time	Freq.	Percent	Cum.
68	1,375	4.82	4.82
69	1,232	4.32	9.14
70	1,686	5.91	15.05
71	1,851	6.49	21.53
72	1,693	5.93	27.47
73	1,981	6.94	34.41
75	2,141	7.50	41.91
77	2,171	7.61	49.52
78	1,964	6.88	56.40
80	1,847	6.47	62.88
82	2,085	7.31	70.18
83	1,987	6.96	77.15
85	2,085	7.31	84.45
87	2,164	7.58	92.04
88	2,272	7.96	100.00
Total	28,534	100.00	

As a result `panelstat` produced three pieces of information. The first block, under the heading *Basic Descriptives*, provides information on the total number of observations, individuals, and the time span of the panel. It also tries to give you an idea of how close the panel is to a fully balanced one. Since we have 4,711 individuals and we observe data for a total of 21 years, if we had a fully

balanced dataset we would have  $4,711 \times 21 = 98,931$  observations. Since we only have 28,534 observations the level of completeness is  $28,534/98,931 = 28.84\%$ .

The second block of information is simply a tabulation of the number of observations per individual. We can immediately see that there are 547 individuals with only one observation (these are called *singletons*) and that only 86 individuals are observed for 15 years - the largest number of periods an individual is observed.

The final block gives us a simple tabulation of the number observations in each year. To avoid calculation of the tables of “Basic Descriptives” we can use the `nosum` option.

### 3 General Options

#### The *excel* option

When working with `panelstat` you may want to save the output of your analyses to an excel spreadsheet. This is done by specifying the name of an output file where you want the results to be stored. For example,

```
panelstat idcode year, excel(myresults)
```

will create an excel spreadsheet that contains all information displayed by the command. By default, the new excel file will be created in Stata’s working directory. You can, however, specify a full path as in

```
panelstat idcode year, excel("C:\mystuff\myresults")
```

If you use other `panelstat` options then all results will be stored in the same excel spreadsheet under different worksheets. To replace an existing spreadsheet you can specify the sub-option `replace` as in

```
panelstat idcode year, excel(myresults, replace)
```

and the sub-option `modify` will let you overwrite an existing spreadsheet.

```
panelstat idcode year, excel(myresults, modify)
```

#### The *force1*, *force2*, and *force3* options

If you try `panelstat` in a dataset that is not a “proper panel” you will obtain an error. This will happen, for example, if there are repeated observations for the unit identifier in the same time period. The options `force1`, `force2`, and `force3` are intended to help you obtain results by ignoring the observations that are causing the problem.

To understand how these options work we will modify the *nlswork* dataset to force an error that will prevent **panelstat** from running. The observations for the first individual in the dataset are

```
. list idcode year if idcode==1
```

	idcode	year
1.	1	70
2.	1	71
3.	1	72
4.	1	73
5.	1	75
6.	1	77
7.	1	78
8.	1	80
9.	1	83
10.	1	85
11.	1	87
12.	1	88

But, we will force now an error in the dataset by replacing the time value of the first 3 observations by 70.

```
. replace year=70 in 1/3
(2 real changes made)
. list idcode year if idcode==1
```

	idcode	year
1.	1	70
2.	1	70
3.	1	70
4.	1	73
5.	1	75
6.	1	77
7.	1	78
8.	1	80
9.	1	83
10.	1	85
11.	1	87
12.	1	88

If you now try to run **panelstat** you will obtain an error because this is not a proper panel. You can force **panelstat** to run by ignoring some observations. There are 3 options:

- **force1** - uses the first observation per repeated value so in this case it ignores the second and third observations.
- **force2** - ignore all repeated values, that is, the first three observations.

- **force3** - ignores all observations for the panel unit so, in this case, all observations for which `idcode==1`.

In our example we only had repeated observations for one individual but if there were multiple individuals with troublesome observations then the same logic would be applied. Note that these options do not modify your dataset – they simply ignore the troubling observations in the ensuing calculations.

### The *cont* option

The **cont** option should be used if there is a time gap that is common to all panel units. Looking at the table produced above with the heading “Number of individuals per time unit” we see that some years are missing. If you use the **cont** option then these years will be ignored in all subsequent calculations. This will make a difference in calculations that make use of lagged values.

### The *forcestata* and *forcesreshape* option

Working with large panels can be extremely slow. Fortunately, there are some user-written tools such as **gtools** written by Mauricio Caceres, **fastreshape** by Michael Droste and **sreshape** by Kenneth L. Simons. By default **panelstat** will check to see if any of these tools is installed in your machine. If it finds the **gtools** package then it will use it. **fastreshape** is also selected by default but if it cannot find it **panelstat** tries to use **sreshape** instead. Only if the user-written commands are not found will it resort to using *Stata* official commands. However, you can change the behavior of **panelstat**. The option **forcestata** will always use *Stata* official commands while the **forcesreshape** will force the use of **sreshape** even if **fastreshape** is installed. Note, however, that the reshape commands are only used with **panelstat**’s options **tabovert** and **pattern**.

## 4 The structure of the panel

There are a set of options that allow us to gain a better understanding of the structure of the panel. These are the **pattern**, **gaps**, **runs**, **vars** and **demog** options. In the following, to prevent redisplaying the basic descriptives for the panel we will use the **nosum** option.

### The *pattern* option

We start with the **pattern** option

```
. use nlswork, clear
(National Longitudinal Survey. Young Women 14-26 years of age in 1968)
. panelstat idcode year, pattern nosum

*****
Top 10 patterns in the data
*****
```

	Pattern	Frequency
1.	1000000000000000	136
2.	0000000000000001	114
3.	0000000000000111	89
4.	0000000000000011	87
5.	1111111111111111	86
6.	000000000011111	61
7.	110000000000000	56
8.	000000111111111	54
9.	000000000001111	54
10.	000000011111111	49

Note: 1 if observation is in the dataset; 0 otherwise

The table that was produced shows the 10 most common patterns in the data. Thus, the most common situation are individuals that were only observed in the first year of the study (136 individuals) while the second most common pattern are individuals observed in the last period (114 individuals). Note that by default we only see the first 10 most common patterns. This behavior can be modified by using the `setmaxpat` option. For example, to report the 25 most common patterns you do

```
panelstat idcode year, pattern setmaxpat(25)
```

## The *gaps* option

A gap is simply a “hole” in the data. Option `gaps` gives you some idea about the gaps in your panel.

```
. panelstat idcode year, gaps nosum

*****
Distribution of the size of the time gaps
*****
```

Size of time gaps	Freq.	Percent	Cum.
1	9,757	75.45	75.45
2	1,014	7.84	83.29
3	737	5.70	88.99
4	493	3.81	92.80
5	108	0.84	93.64
6	240	1.86	95.49
7	109	0.84	96.33
8	85	0.66	96.99

9	115	0.89	97.88
10	42	0.32	98.21
11	86	0.67	98.87
12	34	0.26	99.13
13	44	0.34	99.47
14	21	0.16	99.64
15	17	0.13	99.77
16	12	0.09	99.86
17	8	0.06	99.92
18	8	0.06	99.98
19	2	0.02	100.00
Total	12,932	100.00	
***** Distribution of the number of gaps by individual *****			
Number of gaps per individual	Freq.	Percent	Cum.
0	984	20.89	20.89
1	682	14.48	35.36
2	722	15.33	50.69
3	572	12.14	62.83
4	546	11.59	74.42
5	490	10.40	84.82
6	562	11.93	96.75
7	140	2.97	99.72
8	13	0.28	100.00
Total	4,711	100.00	

The first table that is produced - *Distribution of the size of the time gaps* - gives an idea of how large these gaps are. The most common situation in this data set is a gap of size 1 - which happens 9,757 times. On the other extreme we see that there are 2 gaps of size 19. Since the maximum time dimension of this panel is 21 years this means that two individuals were observed in the first year and again in the last year without having any other observations in between. The second table - *Distribution of the number of gaps by individual* - tells you how many gaps there are per individual. For 984 women there are no gaps - although some of these women may have been observed only once - and, out of the 4,711 women, 13 had a total of 8 gaps!

Notice that there are some years (74, 76, 79, 81, 84 and 86) for which we do not have data. Perhaps there was no data collected in those years, and so this fact may be “inflating” the actual number of gaps. With the `cont` option we can correct this by considering all time periods in the panel as consecutive. If we do this we see that the number of gaps is substantially reduced

```
. panelstat idcode year, gaps nosum cont
*****
Distribution of the size of the time gaps
*****
```

Size of time gaps	Freq.	Percent	Cum.
----------------------	-------	---------	------

1	2,093	55.14	55.14
2	673	17.73	72.87
3	340	8.96	81.82
4	227	5.98	87.80
5	167	4.40	92.20
6	116	3.06	95.26
7	67	1.77	97.02
8	49	1.29	98.31
9	28	0.74	99.05
10	14	0.37	99.42
11	12	0.32	99.74
12	8	0.21	99.95
13	2	0.05	100.00
Total	3,796	100.00	

\*\*\*\*\*  
Distribution of the number of gaps by individual  
\*\*\*\*\*

Number of gaps per individual	Freq.	Percent	Cum.
0	2,171	46.08	46.08
1	1,569	33.31	79.39
2	721	15.30	94.69
3	218	4.63	99.32
4	29	0.62	99.94
5	3	0.06	100.00
Total	4,711	100.00	

If you want you can create variables that capture the number of gaps per panel unit or the size of the largest gap for the panel unit. You do this by adding the option `keepngaps(varname)` or `keepmaxgaps(varname)` where *varname* are names of new variables. For example,

```
. panelstat idcode year, keepngaps(ngap) keepmaxgap(maxgap) nosum cont
```

As a result the variables *ngap* and *maxgap* are added to the dataset. Note that the values for these variables will be different if you do not specify the `cont` option.

## The *runs* option

A *run* is a set of consecutive time periods for which a panel unit is observed. If a panel unit has a run of size 3, this means that this panel unit was observed for three consecutive time periods. Of course, a panel unit may have several *runs* (which meant it also had gaps). Next, we run `panelstat` with the `runs` (as well as the `cont`) option

```
. panelstat idcode year, runs nosum cont
*****
Distribution of complete runs by size
```



\*\*\*\*\*

Length of run	Freq.	Percent	Cum.
1	3,001	35.28	35.28
2	1,635	19.22	54.50
3	1,113	13.08	67.58
4	674	7.92	75.50
5	523	6.15	81.65
6	402	4.73	86.38
7	256	3.01	89.39
8	227	2.67	92.05
9	188	2.21	94.26
10	131	1.54	95.80
11	85	1.00	96.80
12	80	0.94	97.74
13	78	0.92	98.66
14	28	0.33	98.99
15	86	1.01	100.00
Total	8,507	100.00	

We have 3,001 runs of size 1 (some of these are *singletons*) while we have 86 women that are observed for 15 consecutive years (ignoring years for which we do not have any data).

## The *vars* option

The **vars** option produces a table with information for all numeric variables in your dataset indicating how many panel units fall in each of the following categories

- singleton observation with nonmissing value of the variable
- singleton observation with missing value for the variable
- non-singleton with all missing values of the variable
- non-singleton with only one valid value of the variable
- non-singleton with time-invariant values and nonmissing values for the variable
- non-singleton with time-invariant values and missing values for the variable
- non-singleton with time-variant values and nonmissing values for the variable
- non-singleton with time-variant values and missing values for the variable

Running this option in our dataset we obtain

```
. panelstat idcode year, vars nosum cont
```

Distribution of panel units by type of observation for all variables

variable	s_nonmiss	s_missing	allmissing	onevalue	timeinv_nm	t
> imeinv_wm	timevar_nm	timevar_wm				

	idcode	547	0	0	8	4164
>	0	0	0			
	year	547	0	0	8	0
>	0	4164	0			
	birth_yr	547	0	0	8	4164
>	0	0	0			
	age	546	1	0	13	3
>	0	4138	18			
	race	547	0	0	8	4164
>	0	0	0			
<hr/>						
	msp	547	0	0	10	2112
>	4	2036	10			
	nev_mar	547	0	0	10	2895
>	5	1253	9			
	grade	545	2	0	8	4164
>	0	0	0			
	collgrad	547	0	0	8	4164
>	0	0	0			
	not_smsa	547	0	0	8	3422
>	7	734	1			
<hr/>						
	c_city	547	0	0	8	2894
>	5	1262	3			
	south	547	0	0	8	3616
>	6	540	2			
	ind_code	532	15	1	20	1454
>	0	2404	293			
	occ_code	536	11	1	13	1414
>	0	2649	95			
	union	255	292	269	417	527
>	1536	221	1202			
<hr/>						
	wks_ue	499	48	18	225	569
>	793	1026	1541			
	ttl_exp	547	0	0	8	22
>	0	4142	0			
	tenure	535	12	0	25	23
>	0	3757	367			
	hours	546	1	0	12	706
>	0	3392	62			
	wks_work	523	24	1	46	26
>	0	3510	589			
<hr/>						
	ln_wage	547	0	0	8	13
>	0	4151	0			
	maxgap	547	0	0	8	4164
>	0	0	0			
	ngap	547	0	0	8	4164
>	0	0	0			

s\_nonmiss - singleton observation with nonmissing value of the variable  
s\_missing - singleton observation with missing value for the variable  
allmissing - non-singleton with all missing values of the variable  
onevalue - non-singleton with only one valid value of the variable

```

timeinv_nm - non-singleton with time-invariant values and nonmissing values for
> the variable
timeinv_wm - non-singleton with time-invariant values and missing values for th
> e variable
timevar_nm - non-singleton with time-variant values and nonmissing values for t
> he variable
timevar_wm - non-singleton with time-variant values and missing values for the
> variable

```

To create a variable with an indicator showing the above cases for a particular variable you need to use the `wiv` option (see below).

## The *demog* option

Finally, we discuss the `demog` option. This option characterizes the flows of the panel units that occur between consecutive time periods.

```

. panelstat idcode year, demog nosum
*****
Time changes - incumbents, entrants and exits
*****

```

	time	total	inc1	entry	first	reent	inc2	exit	last	reex
68	1375	0	1375	1375	0	851	524	136	388	
69	1232	851	381	381	0	1001	231	79	152	
70	1686	1001	685	476	209	1315	371	93	278	
71	1851	1315	536	381	155	1224	627	156	471	
72	1693	1224	469	331	138	1411	282	97	185	
73	1981	1411	570	257	313	0	1981	132	1849	
75	2141	0	2141	304	1837	0	2141	189	1952	
77	2171	0	2171	275	1896	1625	546	163	383	
78	1964	1625	339	134	205	0	1964	199	1765	
80	1847	0	1847	142	1705	0	1847	143	1704	
82	2085	0	2085	159	1926	1647	438	155	283	
83	1987	1647	340	126	214	0	1987	239	1748	
85	2085	0	2085	147	1938	0	2085	311	1774	
87	2164	0	2164	109	2055	1817	347	347	0	
88	2272	1817	455	114	341	0	2272	2272	0	

```

time - time period
total - total number of individuals at time t
inc1 - number of individuals at t that are also present at t-1
entry - number of individuals at t that are not present at t-1
first - number of individuals at t who show up for the first time at t
reent - number of individuals at t that are reentering at time t
inc2 - number of individuals at t that are also present at t+1
exit - number of individuals at t that are not present at t+1
last - number of individuals at t that are not present at any future time
reexit - number of individuals at t not present at t+1 that appear in later tim

```

```

> es
the following identities hold:
total[t+1]=total[t]-exit[t]+entry[t+1]
inc1=total-entry)
entry=first+reent
inc2=inc1[t+1]
exit=last+reexit

```

If we look at the table we see that in the first year of the panel - 1968 - we had a total of 1,375 individuals. Of those, 851 will be observed again in the following year. Of the remaining 524 which are not observed in 1969, 136 were never again observed in the panel while 388, although not present in 1969, eventually come back in a later year. Moving now to the 1970 row we see that there are 1,686 individuals. 1,001 had been observed in the previous year while 685 are entering the dataset in this year. Of these 685, 476 show up in the data for the first time while 209 had already showed up in a previous year (so they must have gaps). It is no surprise that in 1975 all individuals enter the data - because there was no data in 1974. To account for this we could have specified the `cont` option.

## The *all* option

The `all` option is equivalent to simultaneously selecting the options `pattern`, `gaps`, `runs`, `vars` and `demog`

## 5 Describing your data

`panelstat` offers a set of options that allow you to inspect your variables taking advantage of the panel structure.

### The `statovert` option

This option produces descriptive statistics over time for a list of variables. It is quite simple to use. To obtain year by year statistics for the variable *grade* do

```

. panelstat idcode year, nosum statovert(grade)

```

Descriptive statistics of grade over time

time	total	valid	missing	mean	sd	p25	p50	p75
68	1375	1375	0	11.7847	1.89504	12	12	12
69	1232	1232	0	11.8117	1.85027	12	12	12
70	1686	1686	0	11.8909	1.95588	12	12	12
71	1851	1851	0	11.9773	1.91001	12	12	12
72	1693	1693	0	12.0815	1.95619	12	12	12

73	1981	1981	0	12.2024	1.98884	12	12	12
75	2141	2141	0	12.3349	2.20353	12	12	13
77	2171	2171	0	12.5615	2.42971	12	12	14
78	1964	1964	0	12.7032	2.44511	12	12	14
80	1847	1847	0	12.7991	2.44423	12	12	14
82	2085	2085	0	12.8043	2.47387	12	12	14
83	1987	1987	0	12.9461	2.51593	12	12	15
85	2085	2085	0	13.0432	2.47832	12	12	15
87	2164	2164	0	13.1072	2.45129	12	12	15
88	2272	2270	2	13.0969	2.51577	12	12	15

We could have listed several variables in the argument of `statover`, in which case `panelstat` would create a table for each variable. `statover` also supports the suboption `detail` which, as suggested by the name, provides additional descriptives on the variables (the 1, 5, 95 and 99th quantile of the variable(s)).

## The wiv and wtv options

The `wiv` option provides statistics for a list of selected variables alongside the panel unit dimension. Using this option on the variables *race* and *union*

```
. panelstat idcode year, nosum wiv(race union)
*****
*****
Analyzing variable race within idcode
*****
There are 100.00% nonmissing observations (28534 out of 28534)
For the variable race we have:
  values range from 1 to 3
  547 singleton idcode-observations with non-missing value (11.61%)
  0 singleton idcode-observations with missing value ( 0.00%)
  0 non-singleton idcode-observations with all values missing ( 0.00%)
  0 non-singleton idcode-observations with only one valid value ( 0.00%)
  4164 non-singleton idcode-observations with year invariant and non-missing
> values (88.39%)
  0 non-singleton idcode-observations with year invariant and missing values
> ( 0.00%)
  0 non-singleton idcode-observations with year variant and non-missing valu
> es ( 0.00%)
  0 non-singleton idcode-observations with year variant and missing values (
> 0.00%)
Distribution of all observations for race
```

	_wiv_race	Freq.	Percent	Cum.
1	singleton non-missing	547	1.92	1.92
5	time invariant with non-missing	27,987	98.08	100.00
Total		28,534	100.00	

```
*****
*****
Analyzing variable union within idcode
*****
```

```

There are 67.42% nonmissing observations (19238 out of 28534)
For the variable union we have:
  values range from 0 to 1
  255 singleton idcode-observations with non-missing value ( 5.41%)
  292 singleton idcode-observations with missing value ( 6.20%)
  269 non-singleton idcode-observations with all values missing ( 5.71%)
  409 non-singleton idcode-observations with only one valid value ( 8.68%)
  527 non-singleton idcode-observations with year invariant and non-missing
> values (11.19%)
  1536 non-singleton idcode-observations with year invariant and missing val
> ues (32.60%)
  221 non-singleton idcode-observations with year variant and non-missing va
> lues ( 4.69%)
  1202 non-singleton idcode-observations with year variant and missing value
> s (25.51%)
Distribution of all observations for union

```

_wiv_union	Freq.	Percent	Cum.
1 singleton non-missing	255	0.89	0.89
2 singleton missing	292	1.02	1.92
3 all values missing	745	2.61	4.53
4 one non-missing value	1,310	4.59	9.12
5 time invariant with non-missing	1,986	6.96	16.08
6 time invariant with missing	12,116	42.46	58.54
7 time variant with non-missing	990	3.47	62.01
8 time variant with missing	10,840	37.99	100.00
Total	28,534	100.00	

we obtain a piece of information for each variable. At the top we have information at the level of the individual followed by a table with information at the observation level. Inspecting the output produced by `panelstat` we can see that there are no missing observations for *race*. The panel has 547 singletons (individuals observed only once) but there are no missing values of *race* for the singletons “547 singleton idcode-observations with non-missing value (11.61%)”. The remaining 4,164 individual level observations for *race* are all (as expected) time-invariant and no missing values are recorded “4164 non-singleton idcode-observations with year invariant and non-missing values (88.39%)”. Singletons represent close to 2% of all observations. The situation is more diverse for the *union* variable. A little less than a third of the observations are missing for the *union* variable. The majority of the singleton observations have missing values for this variable (292 out of 547). There are 269 individuals, observed more than once, that have all observations for *union* missing “269 non-singleton idcode-observations with all values missing ( 5.71%)” On the other hand, for 409 individuals there is only one value for *union* recorded with the other values missing for the variable. We can see that 527 women have no missing values *and* have maintained the same status over time while 1,536 also show consistent values over time but have some missing values. A smaller number, 221 women, have no missing values in the variable *union* but change status over time. Finally, 1,202 participants see their union status altered over time but have some missing observations. The distribution in terms of observations shows that the largest proportion of the observations (42.5%) are “time-invariant with missing”. We could retain in

our data a variable indicating the situation of each observation. Running the command with the suboption **keep** as in

```
panelstat idcode year, nosum wiv(race union, keep)
```

would add to the data two variables – `__wiv__race` and `__wiv__union` – to identify the type assigned to each observation.

A less used option is the **wtv** option. It does exactly the same as **wiv** but exchanges the roles of the panel unit and time variable.

## The tabovert option

This is an option to be used with categorical variables. It simply tabulates the frequency of each category over time. Say, we wanted to understand how the variable *union* had changed over time. To avoid working with a wide table we restrict the analysis to the period 1968-1977.

```
. panelstat idcode year if year<78, nosum tabovert(union)
```

Tabulation of union over time

union	n68	n69	n70	n71	n72	n73	n75	n77
0	.	.	620	701	984	817	.	1760
1	.	.	178	227	260	238	.	400
.	1375	1232	888	923	449	926	2141	11

This way we obtain a distribution over time for all values of the variable *union* including missing values.

## The demoby option

This is also an option to be used only with categorical variables. It can be seen as an extension of the previous option. This option tries to provide some idea about the movements over time of the panel units across the categories of a given variable. Thus, if we run the option on the *union* variable we obtain

```
. panelstat idcode year, nosum demoby(union)
```

```
*****
Decomposition of changes across union over time
*****
```

time	total	first	last	sing	stay	mover	fmover	rmover
70	798	798	68	68	0	0	0	0
71	928	457	76	49	429	42	42	0
72	1244	477	123	60	660	107	92	15
73	1055	244	121	36	719	92	66	26
77	2160	933	250	125	1006	221	170	51

78	1351	140	133	25	1076	135	96	39
80	1686	233	114	20	1171	282	207	75
82	2083	257	179	34	1513	313	151	162
83	1805	134	221	25	1465	206	83	123
85	2079	212	347	50	1589	278	136	142
87	2161	153	630	60	1699	309	126	183
88	1888	112	1888	112	1562	214	87	127

Note: missing values of union are discarded for the analysis (to include specif  
> y missing option)  
time - time period  
total - total number of individuals at time t (total=firs+stay+mover)  
first - number of individuals at t that show up for the first time  
last - number of individuals at t that show up for the last time  
singleton - number of individuals at t that show up only that time (singletons)  
stayer - number of individuals at t that were present at the same category of u  
> nion since their last observation  
mover - number of individuals at t that were present at a different category of  
> union since their last observation (mover=fmover+rmover)  
fmover - number of movers at t that are for the first time at that category of  
> union  
rmover - number of movers at t that are returning to a category of union

The above table shows us that in 1970, the first year for which there are nonmissing values of *union*, of the 798 reported values, 68 are for women that only reported values for *union* once. If we move to the row corresponding to the year 1972, we can see that, out of 1,244 valid values reported for the variable *union*, 477 are for women that are reporting their *union* status for the first time, and for 123 women these are the last reported values of *union* status. But there is more information. 660 women maintain the same union status as in the previous year, while 107 have changed their union status. Of those that changed union status, 92 are new to their present category while the other 15 are now returning to a category in which they have been in the past. The suboption *keep* will create a variable identifying each situation at the level of the observation. Thus, if we run

```
panelstat idcode year, nosum demoby(union, keep)
```

the variable *\_demoby\_union* is created and it will contain an indicator of all possible cases. To understand how this variable is coded consider the list of values for this variable for the first individual in the dataset:

```
. list idcode year union _demoby_union if idcode==1
```

	idcode	year	union	_demob_n
1.	1	70	.	.
2.	1	71	.	.
3.	1	72	1	1 first
4.	1	73	.	.
5.	1	75	.	.
6.	1	77	0	3 fmover



7.	1	78	.	.
8.	1	80	1	4 rmover
9.	1	83	1	2 stayer
10.	1	85	1	2 stayer
11.	1	87	1	2 stayer
12.	1	88	1	2 stayer

By default **demoby** does not report information for periods when all values of the variable are missing. That behavior can be overrun with the suboption **missing**

## The flows option

The option **flows** decomposes the change in the stock of each variable between consecutive periods. Thus, for each time period, it identifies the changes that result from panel units that already exist (incumbent), and those that enter and exit. To illustrate the use of the command we consider the variable *hours* which contains the amount of hours worked by each woman.

```
. panelstat idcode year, nosum flows(hours)
*****
Time flows for variable hours
*****
```

	time	hours	chg	c_inc	c_exp	c_cont	c_ent	c_exit	c_incl	c
>	68	51319	.	0	0	0	51319	.	0	
>	69	46713	-4606	536	2096	-1560	14003	-19112	0	
>	70	61858	15145	300	1693	-1393	23315	-8512	42	
>	71	67853	5995	1263	3848	-2585	17721	-12989	0	
>	72	61376	-6477	155	3147	-2992	15765	-22376	0	
>	73	71490	10114	764	2924	-2160	19065	-9727	12	
>	75	78213	6723	0	0	0	78213	-71490	0	
>	77	78248	35	0	0	0	78248	-78213	0	
>	78	70705	-7543	-133	2824	-2957	11137	-18220	40	
>	80	66999	-3706	0	0	0	66999	-70705	0	
>	82	74198	7199	0	0	0	74198	-66999	0	

```

> | 83 70946 -3252 -281 2686 -2967 11645 -14759 400
> | -257 |
> | 85 76396 5450 0 0 0 76396 -70946 0
> | 0 |
> | 87 79711 3315 0 0 0 79711 -76396 0
> | 0 |
> | 88 84716 5005 842 3758 -2916 16158 -12232 237
> | 0 |

```

---

```

Notes:
hours - total sum of hours at time t
chg - sum of hours at t minus t-1
c_inc - changes from individuals present at t and at t-1 with valid values of h
> ours
  of which:
    c_exp - positive changes (expansions) from individuals present at t and at
> t-1
    c_cont - negative changes (contractions) from individuals present at t and
> at t-1
c_entry - change resulting from entry (present at t but not at t-1)
c_exit - change resulting from exits (present at t-1 but not at t)
c_inc1 - change from individuals present at t and t-1 but with missing data at
> t-1
c_inc2 - change from individuals present at t and t-1 but with missing data at
> t
hours[t]=hours[t-1]+chg, chg=c_inc+c_entry+c_exit+c_inc1+c_inc2, c_inc=c_exp+c_
> cont

```

Consider the second row of the table that was produced. It shows that in 1969 the total amount of hours reported was 46,713 and that represented a decrease of 4,606 when compared with the previous year. Incumbents (those women that were present in 1969 and also in the previous year) increased their working hours by 536. The change in hours worked by incumbents can be decomposed in changes by women that reported an increase in hours worked (2,096 hours) and women that decreased the amount of hours worked (1,560 hours). The next column shows the amount of hours added by women that entered the survey in 1969 (were not present in the previous year) - a total contribution of 14,003 while women that exited from 1968 to 1969 contributed with a change of -19,112 hours. We also see that there are no incumbents with missing data in 1968 but there is a reduction of 33 hours which is accounted for by missing data of incumbent(s) who had hours reported in 1968. You can specify several variables and a table will be produced for each variable.

If you want to know how many panel units are behind the numbers produced by the option `flows` you need to specify the *unit* option. That is, you should instead do

```

panelstat idcode year, nosum flows(hours, unit)

```

## 5 Inspecting your data

This set of commands is particularly helpful to spot problems in the data.

### The `rel` and `abs` options

These options report on changes over time for the specified variable. The options are particularly suited for continuous variables. As an example we will use the `rel` option with the variable *hours*

```
. panelstat idcode year, nosum rel(hours) cont
```

```
*****
Relative changes over time for hours (threshold set to 100)
*****
```

<code>_rel_L1_hours</code>	Freq.	Percent	Cum.
1 positive change	4,450	15.60	15.60
2 negative change	3,917	13.73	29.32
3 no change	11,092	38.87	68.20
4 abnormal pos chg	278	0.97	69.17
5 abnormal neg chg	200	0.70	69.87
6 missing	8,597	30.13	100.00
Total	28,534	100.00	

Note: Relative change is calculated with respect to the average of  $x_{\{t\}}$  and  $x_{\{t-1\}}$

The output classifies all relative changes from two consecutive periods for the same panel unit and tabulates the results. It distinguishes between no change, positive and negative change, and abnormal positive, and abnormal negative changes. Abnormal changes are those that exceed a specified threshold value for a relative change (by default 100). You can change the threshold value using the suboption `val`. For example, if we try

```
. panelstat idcode year, nosum rel(hours, val(50)) cont
```

```
*****
Relative changes over time for hours (threshold set to 50)
*****
```

<code>_rel_L1_hours</code>	Freq.	Percent	Cum.
1 positive change	3,792	13.29	13.29
2 negative change	3,411	11.95	25.24
3 no change	11,092	38.87	64.12
4 abnormal pos chg	936	3.28	67.40
5 abnormal neg chg	706	2.47	69.87
6 missing	8,597	30.13	100.00
Total	28,534	100.00	

Note: Relative change is calculated with respect to the average of  $x_{\{t\}}$  and  $x_{\{t-1\}}$

we see that the number of abnormal positive (and negative) changes increases

because we are now classifying as abnormal a relative change exceeding 50%. We should also note that, by default, relative changes are calculated with respect to the average of starting and end point. This behavior can be changed with the `denlag` suboption. If that option is specified then the relative change is calculated with respect to the lagged value. We can also change the number of lags used on the calculation of the relative change. By default that value is 1. But you could specify a different value with the option `lags`. Finally, suboption `keep` creates a variable that stores the classification of type of change attributed to the observations of the variable. The command

```
panelstat idcode year, nosum rel(hours, keep) cont
```

would create a variable named `__rel_L1_hours`.

The option `abs` operates similarly but it classifies the absolute changes. The threshold change to report abnormal changes is 10 and may be changed with the suboption `val`. As with the option `rel` you can use a different value for the lag (the default is 1). This is done with suboption `lags`. If using a lag larger than one you may prefer to use differences in differences. In that case you need to specify suboption `dif`. Finally, you can use the `keep` option to retain the variable with the classifications.

## The fromto option

This is another option intended for categorical variables. Basically, for a given variable, it tabulates all combinations of values at two different time periods. The `cont` option is ignored if you specify the `fromto` option because you have to explicitly identify the time periods. Thus, if I wanted to find out the changes in variable *union* from 87 to 88 I would do

```
. panelstat idcode year, nosum fromto(union, from(87) to(88) )
```

Change of union from 87 to 88

union87	union88	_type	n
0	1	4 dif	80
1	0	4 dif	75
0	0	3 same	1072
1	1	3 same	304

The results show that from 1987 to 1988, 80 women became unionized while 75 left the union. 1072 remained unionized while 304 remained unionized. If we add the suboption `missing` then the missing observations of *union* are also accounted for:

```
. panelstat idcode year, nosum fromto(union, from(87) to(88) missing )
```

Change of union from 87 to 88

--	--	--	--

union87	union88	_type	n
0	1	4 dif	80
1	0	4 dif	75
0	0	3 same	1072
1	1	3 same	304
.	0	2 entry	276
.	1	2 entry	81
0	.	1 exit	527
1	.	1 exit	103

The **from** argument is required and it must contain a valid value for the time variable, while the **to** argument may be omitted. If that is the case then it is assumed that the change is for the following time period. Thus, the same results would obtain had we specified

```
panelstat idcode year, nosum fromto(union, from(87) missing )
```

If you want the displayed table to be sorted from lowest to highest frequency you can use the **ascend** option while the **descend** option does the reverse. **fromto** also supports two other suboptions – **save** and **keep**. The **save** will save a *Stata* file with the results shown in the table, while the **keep** option adds a variable to the dataset that identifies for each observation the following situations:

- 0 not flagged - observation was not considered
- 1 exit - missing value of the variable at *to*
- 2 entry - missing value of the variable at *from*
- 3 same - values are the same
- 4 dif - values are different

In this case the variable name would be `__ft_union_87_88`.

## The return option

This is another option intended to look at the change within panel unit of the values of a specific variable. It considers three time periods say  $x_t$ ,  $x_{t+1}$ , and  $x_{t+2}$ . It checks whether the value of  $x_t$  and  $x_{t+2}$  are identical but  $x_{t+1}$  differs. If we try

```
. panelstat idcode year, nosum return(union, from(71) middle(72) to(73) )
```

Return of union from 71 to 73

union71	union72	union73	n
0	1	0	12
1	0	1	4

then we are requesting `panelstat` to identify observations of union for which values of 71 and 73 are identical with differing values in 72. We see that there are 12 women with a “0 1 0” sequence and another 4 with a “1 0 1” sequence. These may signal coding errors. Since we have three consecutive time values we can simply code

```
panelstat idcode year, nosum return(union, from(71))
```

to obtain the same result. But remember that the arguments for the `from`, `middle` and `to` suboptions must be valid time values. The suboption `save` will save the table of results to a Stata file. You can also use the `keep` suboption to add a variable to the dataset that identifies the flagged observations. In this case the variable would be named `_ret_union_71_73`. The `return` option can also be used with continuous variables. In that case you need to specify the `within` suboption. With this suboption the command will look for cases where  $x_{t+1}$  is outside the interval  $[x_t * (1 - a\%), x_t * (1 + a\%)]$  where  $a$  is the value used in the argument of the `within` suboption. To illustrate let us check for abnormal changes in wages for 72.

```
. panelstat idcode year, nosum return(ln_wage, from(71) within(50))
```

Return of ln\_wage from 71 to 73

ln_wa_71	ln_wa_72	ln_wa_73	n
.8007324	1.292054	.9011876	1
.8462984	1.309547	1.219973	1
.8462984	1.503363	1.231429	1
.9416081	1.454573	.9845695	1
.9830538	1.483829	1.321042	1
1.023876	1.655147	1.480504	1
1.115543	.5257036	1.166891	1
1.206198	.2506002	1.079615	1
1.217862	1.860038	1.394327	1
1.316302	.28339	1.072743	1
1.328725	.3226964	1.525354	1
1.376927	.5522707	.9437474	1
1.434085	2.249825	2.148098	1
1.472237	.3903617	.8322071	1
1.49388	2.370864	1.503363	1
1.790204	.3609478	1.283395	1
1.835839	3.02527	2.707336	1

The above table gives a list of all cases where *ln\_wage* in 72 is outside a 50% percent interval constructed around the 71 value while the 73 value is within that interval.

## The trans option

Again this is an option meant to help you identify potential problems in the data. It is meant to be used with categorical variables. The idea is simple. For all units in each category of the variable in year  $t$  it calculates the share of those same units that came from each different category of the variable at  $t-1$ . We call this the *transition probabilities*. Thus, if in a given year a panel unit has a transition probability of 100% it means that all individuals that belong to the same category at year  $t$  also belonged to the same category at  $t-1$  (but the categories in  $t$  and  $t-1$  need not be the same). Likewise, a transition probability of 10% means that 10% of individuals in a given category at year  $t$  came from the same category where they were classified in  $t-1$ . The results are presented in a table that shows, for each time period, the number of panel units grouped into 4 categories of the *transition probabilities*. To exemplify, let us apply this option to the variable age

```
. panelstat idcode year, nosum trans(age, keep)

*****
Distribution of transition probabilities (t-1 to t) for classes of age
*****
```

Time	Distribution of probabilities				Total
	p<25	25<=p<75	75<=p<100	p=100	
69	48	0	765	31	844
70	48	0	936	17	1,001
71	29	0	1,219	67	1,315
72	41	0	1,071	112	1,224
73	58	0	1,319	34	1,411
78	90	2	1,426	107	1,625
83	71	0	1,536	40	1,647
88	1	1,763	41	0	1,805
Total	386	1,765	8,313	408	10,872

The option `keep` adds the variable `_trans_age` where the transition probabilities are stored. The results are summarized in the table. If we look at the last year, 1988, we can see that we were able to compute transition probabilities for 1,805 individuals (these were individuals with valid values for age in 87 and 88). One individual has a transition probability lower than 25%. Digging into the data we find that the idcode for this individual is

```
. list idcode if _trans_age<25 & year==88
```

	idcode
19224.	3462

and if we list the observations for this individual we obtain

```
. list idcode year age if idcode==3462
```

--	--	--

	idcode	year	age
19218.	3462	75	20
19219.	3462	78	23
19220.	3462	82	27
19221.	3462	83	28
19222.	3462	85	30
19223.	3462	87	32
19224.	3462	88	34

This subject was flagged because it was aged 32 in 1987 and 34 in 1988. It was the only individual out of 42 with valid age values in 1987 and 1988 that moved from age 32 in 1987 to age 34 in 1988 (a “transition probability” of  $1/42=2.38\%$ ).

The suboption `missing` also accounts for transitions from missing to a valid value of the variable. In that case we would have to specify

```
panelstat idcode year, nosum trans(age, miss)
```

We can use the suboptions `low` and `upper` to define the threshold levels used in the table that is displayed. If, say, we want to find out how many women, each year, had transition probabilities below 1% and above 99% we could write:

```
. panelstat idcode year, nosum trans(age, low(1) upper(99))

*****
Distribution of transition probabilities (t-1 to t) for classes of age
*****
```

Time	Distribution of probabilities				Total
	p<1	1<=p<99	99<=p<100	p=100	
69	0	813	0	31	844
70	1	983	0	17	1,001
71	3	1,138	107	67	1,315
72	4	1,108	0	112	1,224
73	1	1,376	0	34	1,411
78	3	1,515	0	107	1,625
83	3	1,604	0	40	1,647
88	0	1,805	0	0	1,805
Total	15	10,342	107	408	10,872

## The `quantr` option

The `quantr` option is also intended to help find problems in the data. However, it is intended for use with continuous data. As an example, we use this option with the `ln_wage` variable. For each year, we compute the 25th, 50th and 75th percentile. Next we look, for consecutive years, how individuals move between these percentiles. Thus, if we do,

```
. panelstat idcode year, nosum quantr(ln_wage)

*****
```



changes (t-1 to t) in the quartiles of ln\_wage

\*\*\*\*\*

		Distribution of quantile changes					
>	Time	1to1	1to2	1to3	2to1	2to2	2to3
	3to1	3to2	3to3	Total			
	69	106	63	16	73	304	56
>	7	49	177	851			
	70	174	135	13	51	435	102
>	8	55	237	1,210			
	71	185	145	21	103	524	105
>	13	76	298	1,470			
	72	172	129	13	108	504	90
>	8	58	280	1,362			
	73	269	150	15	114	657	92
>	14	68	345	1,724			
	75	261	213	29	167	597	148
>	31	102	289	1,837			
	77	279	155	19	178	680	101
>	24	129	331	1,896			
	78	307	120	17	149	679	77
>	12	119	350	1,830			
	80	264	146	11	156	608	107
>	18	99	296	1,705			
	82	339	175	15	130	677	134
>	21	108	327	1,926			
	83	328	89	10	133	733	82
>	19	96	371	1,861			
	85	318	141	15	154	715	93
>	13	135	354	1,938			
	87	347	159	7	146	782	129
>	18	87	380	2,055			
	88	383	147	16	115	837	130
>	29	104	397	2,158			
Total		3,732	1,967	217	1,777	8,732	1,446
>	235	1,285	4,432	23,823			

Notes:

quartile 1 defined as values below 25

quartile 2 defined as values above 25 and below 75

quartile 3 defined as values above 75

we can find the year-to-year movements across the percentiles of *ln\_wage*. Of the 851 individuals in 1969 (those that had non missing wage values in that and the previous year) 106 moved from the first quartile in 1968 to the first quartile in 1969 (1to1 column). More relevant are probably the 16 individuals that had a wage above the 3rd quartile in 1969 but a wage below the 1st quartile in 1968 (1to3 column) If we want we can show the table in terms of shares. For this we use the suboption `rel` as in,

```
. panelstat idcode year, nosum quantr(ln_wage, rel)
```

\*\*\*\*\*

changes (t-1 to t) in the quartiles of ln\_wage

\*\*\*\*\*

	Distribution of quantile changes					

	Time	1to1	1to2	1to3	2to1	2to2	2to3
>	3to1	3to2	3to3	Total			
	69	12.46	7.40	1.88	8.58	35.72	6.58
>	0.82	5.76	20.80	100.00			
	70	14.38	11.16	1.07	4.21	35.95	8.43
>	0.66	4.55	19.59	100.00			
	71	12.59	9.86	1.43	7.01	35.65	7.14
>	0.88	5.17	20.27	100.00			
	72	12.63	9.47	0.95	7.93	37.00	6.61
>	0.59	4.26	20.56	100.00			
	73	15.60	8.70	0.87	6.61	38.11	5.34
>	0.81	3.94	20.01	100.00			
	75	14.21	11.59	1.58	9.09	32.50	8.06
>	1.69	5.55	15.73	100.00			
	77	14.72	8.18	1.00	9.39	35.86	5.33
>	1.27	6.80	17.46	100.00			
	78	16.78	6.56	0.93	8.14	37.10	4.21
>	0.66	6.50	19.13	100.00			
	80	15.48	8.56	0.65	9.15	35.66	6.28
>	1.06	5.81	17.36	100.00			
	82	17.60	9.09	0.78	6.75	35.15	6.96
>	1.09	5.61	16.98	100.00			
	83	17.62	4.78	0.54	7.15	39.39	4.41
>	1.02	5.16	19.94	100.00			
	85	16.41	7.28	0.77	7.95	36.89	4.80
>	0.67	6.97	18.27	100.00			
	87	16.89	7.74	0.34	7.10	38.05	6.28
>	0.88	4.23	18.49	100.00			
	88	17.75	6.81	0.74	5.33	38.79	6.02
>	1.34	4.82	18.40	100.00			
	Total	15.67	8.26	0.91	7.46	36.65	6.07
>	0.99	5.39	18.60	100.00			

Notes:

quartile 1 defined as values below 25

quartile 2 defined as values above 25 and below 75

quartile 3 defined as values above 75

We can, if we want, redefine the cut-off percentiles used to define the quartiles.

For this we use the suboptions `low` and `upper`, as in

```
. panelstat idcode year, nosum quantr(ln_wage, low(10) upper(90))
```

\*\*\*\*\*  
changes (t-1 to t) in the quartiles of ln\_wage  
\*\*\*\*\*

	Time	1to1	1to2	1to3	Distribution of quantile changes		
>	3to1	3to2	3to3	Total	2to1	2to2	2to3
	69	23	38	0	32	632	35
>	2	24	65	851			
	70	48	74	0	37	886	48
>	1	30	86	1,210			
	71	55	86	3	62	1,053	56
>	3	46	106	1,470			

	72	48	84	3	58	983	58
>	1	42	85	1,362			
	73	77	98	3	79	1,259	43
>	3	32	130	1,724			
	75	49	145	3	121	1,265	82
>	6	67	99	1,837			
	77	69	101	1	117	1,343	62
>	7	81	115	1,896			
	78	67	89	6	121	1,317	57
>	1	62	110	1,830			
	80	73	87	3	172	1,144	67
>	4	62	93	1,705			
	82	106	178	4	100	1,273	81
>	2	70	112	1,926			
	83	78	99	2	99	1,340	48
>	4	63	128	1,861			
	85	76	103	4	113	1,388	66
>	1	69	118	1,938			
	87	94	98	1	108	1,483	75
>	2	62	132	2,055			
	88	120	107	8	98	1,534	82
>	3	79	127	2,158			
<hr/>							
	Total	983	1,387	41	1,317	16,900	860
>	40	789	1,506	23,823			

Notes:  
 quartile 1 defined as values below 10  
 quartile 2 defined as values above 10 and below 90  
 quartile 3 defined as values above 90

Now, quartile 1 would correspond to all individuals with wages up to the 10th percentile, while the quartile 3 would correspond to all individuals with wages above the 90th quartile.

The tables created by the `quantr` option ignore the transitions that originate from missing values. For example, the values reported in 1969 ignore individuals that had missing wage data in 1968. We can report these values by adding the `missing` option as in

```
panelstat idcode year, nosum trans(ln_wage, missing)
```

Finally, as with several other options, we may add a variable to the data set that contains information about the case that applies to each observation. This is done using the suboption `keep`.

## Miscellaneous

### The `checkid` option

This option is used if you have an alternative identifier for the panel unit var and want to find out if how close that variable is to your known panel unit identifier. For example, suppose that in this dataset, besides the `idcode` you had available

another variable – say the tax id number (*taxid*). To understand how close *taxid* is to a panel unit identifier you can compare it to *idcode*. This is done by running the `checkid` option as in

```
panelstat idcode year, checkid(taxid)
```

This option will produce a table with the number of panel units that fit in each of the following cases:

- 1 - 1:1 ids coincide - idcode and taxid coincide
- 2 - 1:m multiple values of taxid - one idcode corresponds to multiple values of taxid
- 3 m:1 multiple values of id - there are multiple idcodes with the same taxid
- 4 m:m multiple values of taxid and id - there are multiple idcodes mixed with multiple taxid
- 5 1:. all values missing for taxid " - one idcode with all values missing for taxcode
- 6 1:1 unique values of taxid with missing - one idcode with unique values of taxid but with missing values
- 7 1:.m multiple values of taxid with missing - one idcode with multiple values of taxid and missing
- 8 m:. multiple values of id with missing - multiple values of idcode with multiple taxids and missing

You can create a variable that stores each one of these cases using the suboption `keep`. Note that to run this option you need to previously install `group2hdfe` a Stata user-written command available at SSC. If not already installed simply type

```
ssc install group2hdfe
```

at the Stata prompt.

## 6 Acknowledgements

I am grateful for all the comments, suggestions, and bug corrections by the BPlim staff, particularly Marta Silva and Emma Zhao. Other researchers, such as Ana Rute Cardoso, have also provided valuable suggestions.